

KPIs, variables, grain, chart critique, bounded data stories

Quantitative Methods for Business Analysis. The essentials of every section in one short read: ideas, worked examples and checks.

≈ 19 min read

Download the PDF

Full edition

1 · READ

Open a section, read its central idea, then its worked example.

2 · APPLY

Try the guided practice and the lab task linked in each section.

3 · CHECK

Answer the study check in the writing space without looking back.

Your study route

2.1

Decision frame, KPIs, targets and guardrails

8 blocks · slides 9–19 · Part A · Part B

2.2

Reproducible KPI definitions, aggregation, and center and position

8 blocks · slides 20–38 · Part C · Part D · Part L

2.3

Variables, measurement scales, observations and grain

9 blocks · slides 39–54 · Part E · Part F · Part K

2.4

Honest charts, bounded stories and validation

8 blocks · slides 55–70 · Part G · Part H · Part I · Part J · Part M · Part N

Contents

2.1 Decision frame, KPIs, targets and guardrails	4
Start with the decision	4
Give every number one role	5
Worked example: MetroMeal's latest week	5
A short workflow	5
Deeper explanation: when the metric becomes the goal	6
Common misconceptions	6
Limits of the evidence	6
Study actions	7
2.2 Reproducible KPI definitions, aggregation, and center and position	8
Write a reproducible specification	8
Combine periods by counts, not by averaging rates	9
Worked example: how late were the late deliveries?	9
A short workflow	10
Deeper explanation: comparing the four weeks	10
Common misconceptions	11
Limits of the evidence	11
Study actions	11
2.3 Variables, measurement scales, observations and grain	13
Columns, rows and two classifications	13
Worked example: one MetroMeal observation	14
Grain and keys	14
Worked example: the join fan trap	14
A short workflow	15
Deeper explanation: what a complete file still cannot show	15
Common misconceptions	15
Limits of the evidence	16
Study actions	16
2.4 Honest charts, bounded stories and validation	17
Choose the display for the question	18
Worked example: repairing Chart A	18

Write a bounded story	18
Validate independently	19
Deeper explanation: one set of facts, several audiences	19
Common misconceptions	20
Limits of the evidence	20
Study actions	20

2.1 · Decision frame, KPIs, targets and guardrails

Learning purpose: Turn a vague request into a decision question, then calculate and read a KPI against its target and guardrail.

CENTRAL IDEA

A number becomes evidence only when you know which decision it informs, which role it plays and what it cannot show.

IN CLASS

[Slides 9–19](#)

Decision-First Analysis 10–11

KPI, Target, and Guardrail 12–14

MetroMeal's Bounded Question 15–16

Retrieval Check 18

PRACTISE IT IN THE LAB

- [Part A: What decision can the KPI inform?](#) · 3 min
- [Part B: Is MetroMeal's latest week below target?](#) · 3 min · SPSS Block H1: decision, KPI, target, and guardrail

About 4 min read · PDF pages 4–7 · [Full explanation \(Full Toolkit\)](#)

Start with the decision

Decision-first analysis defines the decision before touching the data. Name five things: the **stakeholder** (who will act), the **decision** (the realistic choice), the **action horizon** (when the choice must be made), the **outcome** that matters and the **evidence** needed. A useful prompt is **Who, What, and How**: who will use the communication, what they need to know or do, and how the evidence will reach them.

Keep two periods apart. The **observed window** is when the records were collected. The **action horizon** is when the decision takes effect. Data from the past window cannot by themselves describe the future horizon.

A **KPI**, or key performance indicator, is a quantified measure used to monitor an objective and support a decision. It informs the decision; it is not the decision.

Give every number one role

Role	Meaning	MetroMeal
KPI	The measure being monitored	Weekly on-time delivery rate
Target	A management benchmark, chosen rather than calculated	At least 85.0% on time (a minimum floor)
Guardrail	A second measure that must not deteriorate while the KPI is pursued	Complaint rate at most 3.0% (a maximum ceiling)

A **count** is a frequency, such as 392 on-time deliveries. A **proportion** divides that count by the eligible total, giving a value from 0 to 1, often shown as a percentage. A **target gap** subtracts the target from the observed rate. When both are percentages, the direct difference is in **percentage points**. **Percent change** divides the difference by a baseline and describes relative size.

Worked example: MetroMeal's latest week

The regional operations director must decide whether to investigate and potentially adjust weekend staffing before the next scheduling cycle. The data hold one row per eligible completed delivery, August 3 through August 30, 2026: 2,000 rows, 500 per week. Constraints are a fixed staffing budget, customer privacy, incomplete causal information and protecting complaint performance.

For the week of August 24:

```
Reconcile:      392 on time + 108 late = 500 eligible
On-time rate:   392 / 500 = 0.784 = 78.4%
Target gap:     78.4% - 85.0% = -6.6 percentage points
Percent change: -6.6 / 85.0 = -7.76% (relative to the target level)
Complaint rate: 14 / 500 = 2.8%, within the 3.0% guardrail
```

The week missed its KPI target and passed its guardrail. Both statements are true at once, and neither cancels the other. Write the unit, population and week into the sentence: "78.4% of 500 eligible completed deliveries in the week of August 24 were on time."

A short workflow

1. Rewrite the request as a decision question naming stakeholder, decision, window, outcome, comparison, target and guardrail.

2. Draft a provisional claim; mark words that need data and words that imply a cause.
3. Reconcile the counts before dividing.
4. Calculate the KPI and the guardrail with the same eligible denominator.
5. Report the gap in percentage points, and name the baseline if you also give a percent change.

GO DEEPER Deeper explanation: when the metric becomes the goal

Targets carry different meanings. MetroMeal's 85.0% is a minimum **floor**, so a miss should trigger an operational review; the 3.0% complaint limit is a maximum **ceiling**, so a breach should trigger customer-impact checks. A benchmark without an owner, defined population, effective period or breach response is only a label on a chart.

Once a measure becomes a target, people may improve the score without improving the service. The session calls this **metric gaming** and links it to Goodhart's law. MetroMeal teams could change eligibility, classify borderline deliveries as on time, mark orders complete before handoff or rush drop-offs. Stable definitions, visible denominators, audits and guardrails reduce the risk. Still ask the session's ethical check: who is excluded, pressured or harmed by this target?

COMMON MISTAKE Common misconceptions

- **"A 6.6% drop."** The gap came from subtraction, so it is 6.6 percentage points. The relative shortfall is 7.76%.
- **"The guardrail passed, so performance is fine."** The two thresholds protect different outcomes. Report both.
- **Complaints over on-time deliveries.** $14 \div 392 = 3.6\%$ would suggest a false breach. Complaints are counted among all 500 eligible deliveries.
- **The chart is the deliverable.** A draft such as "Performance collapsed because staffing was inadequate" pairs an undefined dramatic word with a causal claim.

LIMITS AND CONDITIONS Limits of the evidence

The records are **observational and descriptive**: they show what happened, not why. Weather, restaurant delays, order mix or staffing could be involved, and the case separates none of them. One week shows a level, not a trend. A gap from target is not automatically statistically significant or managerially important. The on-time rate also hides severity, which Section 2.2 addresses. A guardrail can miss harms, such as pressured couriers or customers who never complain.

Study actions

Reflect: The director wants a dramatic chart for a meeting. How would you balance that request with your duty to communicate honestly?

Recall without looking back: Define KPI, target and guardrail. Why is -6.6 reported in percentage points?

Lab connection: Complete Part A (HarborCare, Class_Practice!B6:B17; checks 11.5% and 3.0%) and Part B (MetroMeal, Slide_Practice!B6:B17; checks 78.4%, -6.6 and 2.8%). Run SPSS Block H1 when directed and compare its values.

2.2 · Reproducible KPI definitions, aggregation, and center and position

Learning purpose: Specify a KPI so that others can rebuild it, combine periods correctly, and describe how late the late deliveries were.

CENTRAL IDEA

The same label can hide different rules, weights and experiences, so write the rules down and show more than one summary.

IN CLASS

[Slides 20–38](#)

Reproducible KPI Specification 21–22

Count, Proportion, and Target Gap 23–25

Percentage Points vs Percent Change 26–27

Reproduce the KPI in Excel 365 28

Aggregation from Counts 32–34

PRACTISE IT IN THE LAB

- [Part C: Is the KPI calculation reproducible?](#) · 3 min
- [Part D: What is the four-week on-time rate, built from counts?](#) · 3 min · SPSS Block H2: weekly and aggregate KPI calculations
- [Part L: How did each week compare with the target and the guardrail?](#) · 4 min · SPSS Guided Step 2: weekly KPI, guardrail, aggregate, and late-delivery checks

About 5 min read · PDF pages 8–12 · [Full explanation \(Full Toolkit\)](#)

Write a reproducible specification

A KPI is **reproducible** when another analyst could compute the same number from your definition alone. A **KPI specification** states the name, purpose, unit, eligible population, time window, grain, inclusion and exclusion rules, numerator or aggregation rule, denominator, direction of improvement, target, guardrail and refresh schedule. The **eligible population** is the set of records allowed into the denominator. For MetroMeal's latest week, the denominator is 500 completed deliveries from August 24 through August 30, one row per delivery, with canceled and test orders excluded.

Any rule change creates a new **KPI version**. The session's example: if "on time" moved from 45 to 50 minutes, the rate could rise with no operational improvement. Record the rule, date, reason and approver, and warn users when comparisons break.

Combine periods by counts, not by averaging rates

Count-first aggregation sums the qualifying counts, sums the eligible counts, then divides once, so each delivery counts once.

Week start	Eligible	On time	Rate	Gap (points)	Complaints	Complaint rate
2026-08-03	500	420	84.0%	-1.0	10	2.0%
2026-08-10	500	438	87.6%	+2.6	11	2.2%
2026-08-17	500	425	85.0%	0.0	12	2.4%
2026-08-24	500	392	78.4%	-6.6	14	2.8%
Four weeks	2,000	1,675	83.75%	-1.25	47	2.35%

The average of the four weekly rates is also 83.75%, but only because every week has 500 deliveries. With unequal weeks, the shortcut fails. The lab's example: 90 of 100 (90.0%) and 300 of 400 (75.0%) average to 82.5%, but only 390 of 500 deliveries were on time, 78.0%. A **delivery-weighted** rate answers "what share of all deliveries was on time?"; a **week-weighted** average answers "what did the average week look like?"

Read the table both ways. The pooled 83.75% is correct, yet three strong weeks hide the weak latest week. The director needs the weekly rows.

Worked example: how late were the late deliveries?

The on-time rate counts lateness but not severity. For the 108 late deliveries in the latest week, the session reports a **mean** of 12.4 minutes, a **median** of 7 minutes and a **90th percentile** of 31 minutes. The mean is the total divided by the count. The median is the ordered middle value. A percentile is a position: about 90% of late values fall at or below 31 minutes under Excel's `PERCENTILE.INC` rule.

This edition tabulated the 108 late_minutes values from `orders_story.csv` (latest week, eligible, late_minutes above 0; every promised time is 45 minutes):

Minutes late	Deliveries	Sorted positions
2.0	20	1 to 20
7.0	35	21 to 55
12.0	41	56 to 96
31.0	2	97 to 98
50.0	9	99 to 107
50.2	1	108

Mean: 1,339.2 total minutes / 108 = 12.4 minutes
Median: positions 54 and 55 are both 7, so 7 minutes
90th percentile: position $1 + 0.9 \times 107 = 97.3$; positions 97 and 98 are 31, so 31 minutes

The mean sits 5.4 minutes above the median because ten deliveries were about 50 minutes late. The median describes the typical late experience; the mean describes the average burden; the percentile describes the longer delays. A mean above the median is consistent with a **right-skewed** distribution, but the session treats it only as a reason to inspect the distribution in Session 3.

A short workflow

1. Write the specification before calculating.
2. Build each week's numerator and denominator from the raw rows.
3. Sum counts first for any overall rate, and show the weekly rows beside it.
4. Pair the frequency KPI with a severity summary, stating the subset and percentile convention.

GO DEEPER Deeper explanation: comparing the four weeks

The same late-minute summaries can be computed for every week, as in the lab's optional Topic 4. This edition calculated them from `orders_story.csv` with pandas (late_minutes above 0 within each week; 90th percentile by the PERCENTILE.INC rule):

Week start	Late deliveries	Mean (min)	Median (min)	90th percentile (min)
2026-08-03	80	11.4	11	19
2026-08-10	62	11.6	12	19
2026-08-17	75	11.4	11	18.6
2026-08-24	108	12.4	7	31

Only the latest week shows a mean well above its median and a 90th percentile near half an hour. That contrast tells you where to look next. It does not tell you why those deliveries were late.

COMMON MISTAKE Common misconceptions

- **Averaging rates is always safe.** It works here only because the weeks are equal in size.
- **The mean is the typical delay.** 96 of the 108 late deliveries were 12 minutes late or less, below the 12.4 mean.
- **These summaries describe every delivery.** They describe the 108 late deliveries in one week, not all 500.
- **A definition can change quietly.** Changing the denominator creates a new KPI version.

LIMITS AND CONDITIONS Limits of the evidence

A specification makes a KPI repeatable, not correct: a wrong eligibility flag is repeated faithfully. Three summary numbers cannot show whether long delays clustered on particular days or regions, and they cannot identify a cause. Percentile values vary by software; SPSS may differ slightly from 31. Claims about skew, outliers and shape wait for Session 3.

Study actions

Reflect: Which of these is most useful for a service promise to customers: 392 on-time deliveries, 78.4%, a 7-minute median delay or a 31-minute 90th percentile? Why does the answer depend on the decision?

Recall without looking back: Why does count-first aggregation give each delivery equal weight? Why can the mean and median both be correct yet answer different questions?

Lab connection: Complete Part C (NorthLine, Class_Practice!B21:B35; checks 90.0% and -2.0) and Part D (Slide_Practice!B21:B35; four-week rate 83.75%). In Part L, build KPI_Build rows 7 to 11 and the late-delivery checks in row 18 (108, 12.4, 7, 31, and 0 in G18). Run Block H2 and Guided Step 2 when directed.

2.3 · Variables, measurement scales, observations and grain

Learning purpose: Classify each field correctly and confirm what one row represents before any KPI is calculated.

CENTRAL IDEA

Meaning, not storage, decides the arithmetic, and row grain decides the weights and the denominator.

IN CLASS

[Slides 39–54](#)

Variables and Observations 40–41

Data Type and Measurement Scale 42–43

Row Grain and Supported Claims 45–46

Mixed-Grain Join Risk 47

Visualizing the Join Fan Trap 48–49

PRACTISE IT IN THE LAB

- [Part E: What does one row represent?](#) · 3 min
- [Part F: How many orders hide in four joined rows?](#) · 3 min · SPSS Block H3: item grain, repeated order IDs, and restoration to order grain
- [Part K: What does one row of the MetroMeal file represent?](#) · 4 min · SPSS Guided Step 1: source, candidate key, grain, and eligibility

About 4 min read · PDF pages 13–16 · [Full explanation \(Full Toolkit\)](#)

Columns, rows and two classifications

A **variable** is a recorded characteristic and appears as a column. An **observation** is the full set of values for one member or event and appears as a row. Each field then receives two separate classifications.

Data type: **categorical** (labels, codes, flags) or **numerical**, which is **discrete** when counted and **continuous** when measured. **Measurement scale:** **nominal** (names without order), **ordinal** (ordered, spacing unknown), **interval** (equal gaps, no true zero) or **ratio** (equal gaps and a true zero). Storage format is a third, unrelated property. Jersey number 10 is not twice jersey number 5.

Worked example: one MetroMeal observation

Order 10482 appears in the slides and in `orders_story.csv` : completed Saturday, August 8, 2026, North region, 3 items, promised 45 minutes, delivered in 54, 9 minutes late, satisfaction category 2, not on time, no complaint.

Field	Value	Data type	Scale	Role
order_id	10482	Categorical	Nominal	Identifier (candidate key)
region	North	Categorical	Nominal	Category
satisfaction_category	2	Categorical	Ordinal	Ordered category
item_count	3	Numerical, discrete	Ratio	Measure
delivery_minutes	54.0	Numerical, continuous	Ratio	Measure
on_time_kpi_flag	0	Categorical (binary)	Nominal	Flag

Averaging order_id has no meaning. Satisfaction ratings are ordered but not equally spaced, so report counts and the median; in the file they run 1: 47, 2: 278, 3: 557, 4: 560 and 5: 558 deliveries (counted from `orders_story.csv`). Item count and minutes support sums, means and ratios. A 0/1 flag is categorical, yet summing it counts the 1s and averaging it gives their proportion. That is how the lab counts eligible deliveries and computes rates.

Grain and keys

Grain is the meaning of one row: entity or event, level of detail and time unit. Test it by completing "Each row represents one ...". A **candidate key** is a field that is unique at that grain. MetroMeal's raw file: each row is one eligible completed delivery; 2,000 rows and 2,000 distinct order_id values, so the key is unique. The weekly summary: each row is one week, four rows. The deliveries also come from 1,000 distinct customers, two deliveries each (counted from the file), so a 78.4% delivery rate does not mean 78.4% of customers were served on time.

Worked example: the join fan trap

In a **one-to-many join**, each parent row repeats once per matching child row. The session's illustration: Order A was late with three items; Order B was on time with one item.

Order grain: 1 on time / 2 orders = 50.0%
After the join: A appears 3 times, B once
Row-based rate: 1 on-time row / 4 rows = 25.0%

Nothing about delivery changed; the late order simply gained triple weight. In MetroMeal's file, 1,675 deliveries have one item and 325 have three (from item_count), so a direct join to an item table would create $1,675 + 3 \times 325 = 2,650$ rows. The repair is to aggregate items to one row per order first, then verify key, relationship and row count.

A short workflow

1. Classify every field twice: type and scale.
2. Complete the grain sentence for each table.
3. Compare row count with distinct-key count.
4. Before joining, identify the relationship and aggregate child rows to the grain the KPI needs.
5. Check that the denominator is unchanged after any join.

GO DEEPER Deeper explanation: what a complete file still cannot show

Cross-sectional data describe many entities at one period; **time-series data** follow a measure through ordered periods. MetroMeal's weekly table is a four-point time series, which is too short to establish a stable trend. The raw file may contain every eligible delivery in its window and still not represent future weeks, excluded orders or a broader customer population.

Dates deserve care. They are ordered and support interval calculations under defined conventions, but Excel's numeric storage does not make them ordinary quantities. Choose the grain by the decision: raw delivery rows to validate records, weekly rows to compare weekly performance, and item rows only when the question concerns items, with the population and denominator redefined.

COMMON MISTAKE Common misconceptions

- **"If Excel can average it, it is a number."** Digits can label, rank or measure.
- **Nominal versus discrete.** Scale and type are different questions; item_count is numerical discrete and ratio-scaled.
- **Deleting repeated IDs cleans a joined table.** At item grain the repeats are correct. Decide the grain, then aggregate.

- **Four joined rows are four orders.** Customers who buy more items quietly gain more say in the rate.



LIMITS AND CONDITIONS

Limits of the evidence

A passing key test shows uniqueness, not that every row is real, eligible or complete. Recorded precision does not settle type, and dates need explicit conventions. An aggregated weekly row cannot be taken apart to recover individual deliveries. If a decision concerns item defects, item grain is right, but the population and denominator must be redefined.

Study actions

Reflect: How would the business story change if one row represented an item, a customer, a day or a week?

Recall without looking back: Why is order_id nominal? Why did the joined rate fall from 50.0% to 25.0%?

Lab connection: Complete Part E (BrightCart, Class_Practice!B39:F52; 310 joined rows), Part F (Slide_Practice!B39:F52; 4 rows, 2 orders) and Part K (KPI_Build!B3:B5 each show 2000; decision frame in Guided_Practice!B5). Run Block H3 and Guided Step 1 when directed.

2.4 · Honest charts, bounded stories and validation

Learning purpose: Choose and repair a chart, then write a short story whose every number is checked and whose claims stay inside the evidence.

CENTRAL IDEA

Accurate numbers can still mislead through scale, labels or wording, so the chart and the story must be as bounded as the data.

IN CLASS

Slides 55–70

Match the Display to the Question 56–58

Axes, Baselines, and Statistical Meaning 59–61

Chart Critique and Accessibility Protocol 62–63

Apply the Critique in Excel 365 64

Bounded Chart Interpretation 65–66

PRACTISE IT IN THE LAB

- [Part G: Does the chart communicate honestly?](#) · 3 min
- [Part H: What makes the MetroMeal dashboard chart misleading?](#) · 3 min · SPSS Block H4: four weekly rates and a simple zero-baseline bar chart
- [Part I: Is the data story traceable and bounded?](#) · 3 min
- [Part J: Can every number in the MetroMeal story be traced?](#) · 3 min · SPSS Block H5: reconcile the latest evidence and calculate story values
- [Part M: Does the chart show the checked numbers honestly?](#) · 4 min · SPSS Guided Step 3: weekly on-time chart from checked values
- [Part N: What should the director do next?](#) · 3 min · SPSS Guided Step 4: values to validate against the chart title and bounded story

About 5 min read · PDF pages 17–20 · [Full explanation \(Full Toolkit\)](#)

Choose the display for the question

Viewer's task	Display	Key check
Look up exact values	Table	Units, denominators, period
Compare magnitudes across categories or weeks	Bar chart	Zero baseline, labels, ordering
Follow an ordered sequence	Line chart	Labeled range, all points, target

A plotted weekly rate is a proportion with a denominator; a target line is a benchmark, not a fitted value. Bars are read by length, so they need a zero baseline. Line charts encode position and may use a narrower range when the range and context are clearly labeled. **Clutter** (3D effects, heavy gridlines, redundant labels, excess color) consumes attention without adding evidence. **Accessible** charts use readable text, contrast, direct labels and a cue beyond color, plus a short text description.

Worked example: repairing Chart A

The session's Chart A shows the four weekly rates as 3D bars on a 75% to 90% axis, uses red and green as the only cues, omits denominators and is titled "Delivery performance collapses."

```
Truncated axis: 78.4% bar fills  $(78.4 - 75) / 15 = 22.7\%$  of the height
                87.6% bar fills  $(87.6 - 75) / 15 = 84.0\%$ 
                apparent ratio  $12.6 / 3.4 = 3.7$  times
Zero baseline:   $87.6 / 78.4 = 1.12$ , a difference of about 12%
```

The repair: a 0% to 100% two-dimensional bar chart, direct labels such as "78.4% (n = 500)", a labeled 85.0% target line, a pattern or outline on the latest bar, and the title "Week of August 24 on-time delivery was 78.4%, 6.6 percentage points below target." The data are identical; only encoding and wording change. The repaired chart shows a real shortfall, 9.2 points below the 87.6% week, without implying collapse, a trend or a cause.

Write a bounded story

Exploratory analysis searches for patterns; **explanatory communication** selects checked evidence for a defined audience and decision. A **bounded story** follows five parts: audience and decision; observed result; comparison with target or prior values; bounded action; limitation and next check. The strongest verb sets the evidence standard: "observed", "lower than" and "associated with" fit descriptive data; "predicts" and "caused" need more.

The session's model story: "For the operations director deciding whether to investigate weekend staffing before the next schedule, on-time delivery for August 24 through August 30 was 78.4%, or 392 of 500 eligible deliveries. This was 6.6 percentage points below the 85% target and lower than the prior three observed weeks, which ranged from 84.0% to 87.6%. The complaint rate remained within its 3% guardrail at 2.8%. Review order-level timing by day and relevant staffing records before reallocating resources because four descriptive weeks do not establish the cause or the likely effect of a staffing change."

An exploratory clue from the lab's optional Topic 2 shows why "investigate" fits: in the latest week, 72 of 142 weekend deliveries (50.7%) were on time, against 320 of 358 weekday deliveries (89.4%). Weekend staffing is one hypothesis; weather, demand, routes and order mix may also differ on weekends.

Validate independently

Validation challenges a result by a different route, not by retyping the same formula.

```
Reconcile: 392 + 108 - 500 = 0
Recompute: 392 / 500 = 78.4%; 85.0 - 78.4 = 6.6 points; 14 / 500 = 2.8%
Compare:   chart labels = KPI_Build cells = numbers in the story
```

The traceability chain runs from decision to variables and grain, KPI and target, checked statistic, honest chart, bounded claim, action, limitation and next check. If any link disagrees, stop and resolve it before recommending action.

GO DEEPER Deeper explanation: one set of facts, several audiences

The same evidence can be adapted without changing its facts. An executive version leads with the target miss and the decision to investigate. An analyst version adds the 392 of 500 denominator, the KPI definition, the source checks and the four-week limitation. A live presenter can explain definitions aloud; a circulated document cannot, so it must carry labels, denominators, the time window and the limits on the page. Every MetroMeal display keeps three pieces of context: 500 deliveries per week, the 85.0% target and the August 3 to August 30 window.

A causal claim would need more than description: credible variation in staffing, aligned comparison periods, reliable scheduling records and checks for other explanations. Until then, the defensible action is to investigate.

**COMMON MISTAKE****Common misconceptions**

- **Fixing colors fixes the chart.** A 75.0% baseline still exaggerates the dip.
- **Typed chart values are fine.** Link chart cells to the checked table, or later corrections will not reach the chart.
- **A real miss proves its cause.** It does not. Remove "collapse", "crisis", "significant", "caused" and "will" unless supported.
- **All ticks mean sound reasoning.** Self_Check compares numbers and counts words; it cannot detect a causal claim.

**LIMITS AND CONDITIONS****Limits of the evidence**

An honest chart cannot repair a wrong denominator or grain. Four descriptive weeks cannot establish a trend, a forecast or a cause. The session's forward-check list for Session 3 covers missing completion times, duplicate order identifiers, changed eligibility, unusual delays, and whether mean, median, spread and shape agree. Avoid generic limits such as "More research is needed"; name the specific gap in the evidence.

Study actions

Reflect: Highlighting the latest week focuses attention on a real problem but can create drama. What emphasis and title would you choose, and why?

Recall without looking back: Why must bar charts start at zero when line charts need not? Name the three sentences of a bounded story.

Lab connection: Complete Parts G to J (Class_Practice and Slide_Practice rows 56 to 88; RiverTown checks 90.0%, 90.0%, 60.0%; GreenBlock checks 91.0% and -4.0). In Part M, link Chart_Build!A5:E8 to KPI_Build and build the combo chart with a 0 to 1 axis. In Part N, write three bounded sentences in Response!B7 after every Self_Check row shows ✓. Run Blocks H4 and H5 and Guided Steps 3 and 4 when directed.